

01 Reproduction

Isolated Lab Results

Project: L_VLLM

Tier: TIER_4_INFERENCE_AGENTS

Identity: Upstream vllm-project/vllm @ 1117140edba8 (Apache-2.0)

Why isolation is recorded separately

Results produced locally are not the same claim as results produced after a public launch. Local runs can be shaped by the operator; public runs cannot. Conflating them would let a controlled experiment masquerade as a field result, which is the specific failure the II budget exists to prevent. This folder holds only the isolated, pre-launch condition.

Reproduction contract

A result is admissible here only if all of the following are recorded:

1. The exact commit of this project and of its upstream.
2. The environment (interpreter, platform, installed extras).
3. The command, verbatim.
4. The pass condition, stated before the run.
5. The observed output, unedited.

Current state

Project L_VLLM has no executed lab result recorded at the time of generation. The register below is the template the first run must fill in; it is not a result.

Result	Pass condition	Command	Observed	Status
--------	----------------	---------	----------	--------

---	---	---	---	---
-----	-----	-----	-----	-----

(none yet)	(define before running)	(verbatim)	(unrecorded)	NOT RUN
------------	-------------------------	------------	--------------	---------

After public launch

Kaggle and Hugging Face deployment results belong in a separate record, never here, and must state the model weights, dataset revision, and inference configuration used. Expect these to differ from the isolated figures; where they do, the difference is the finding.